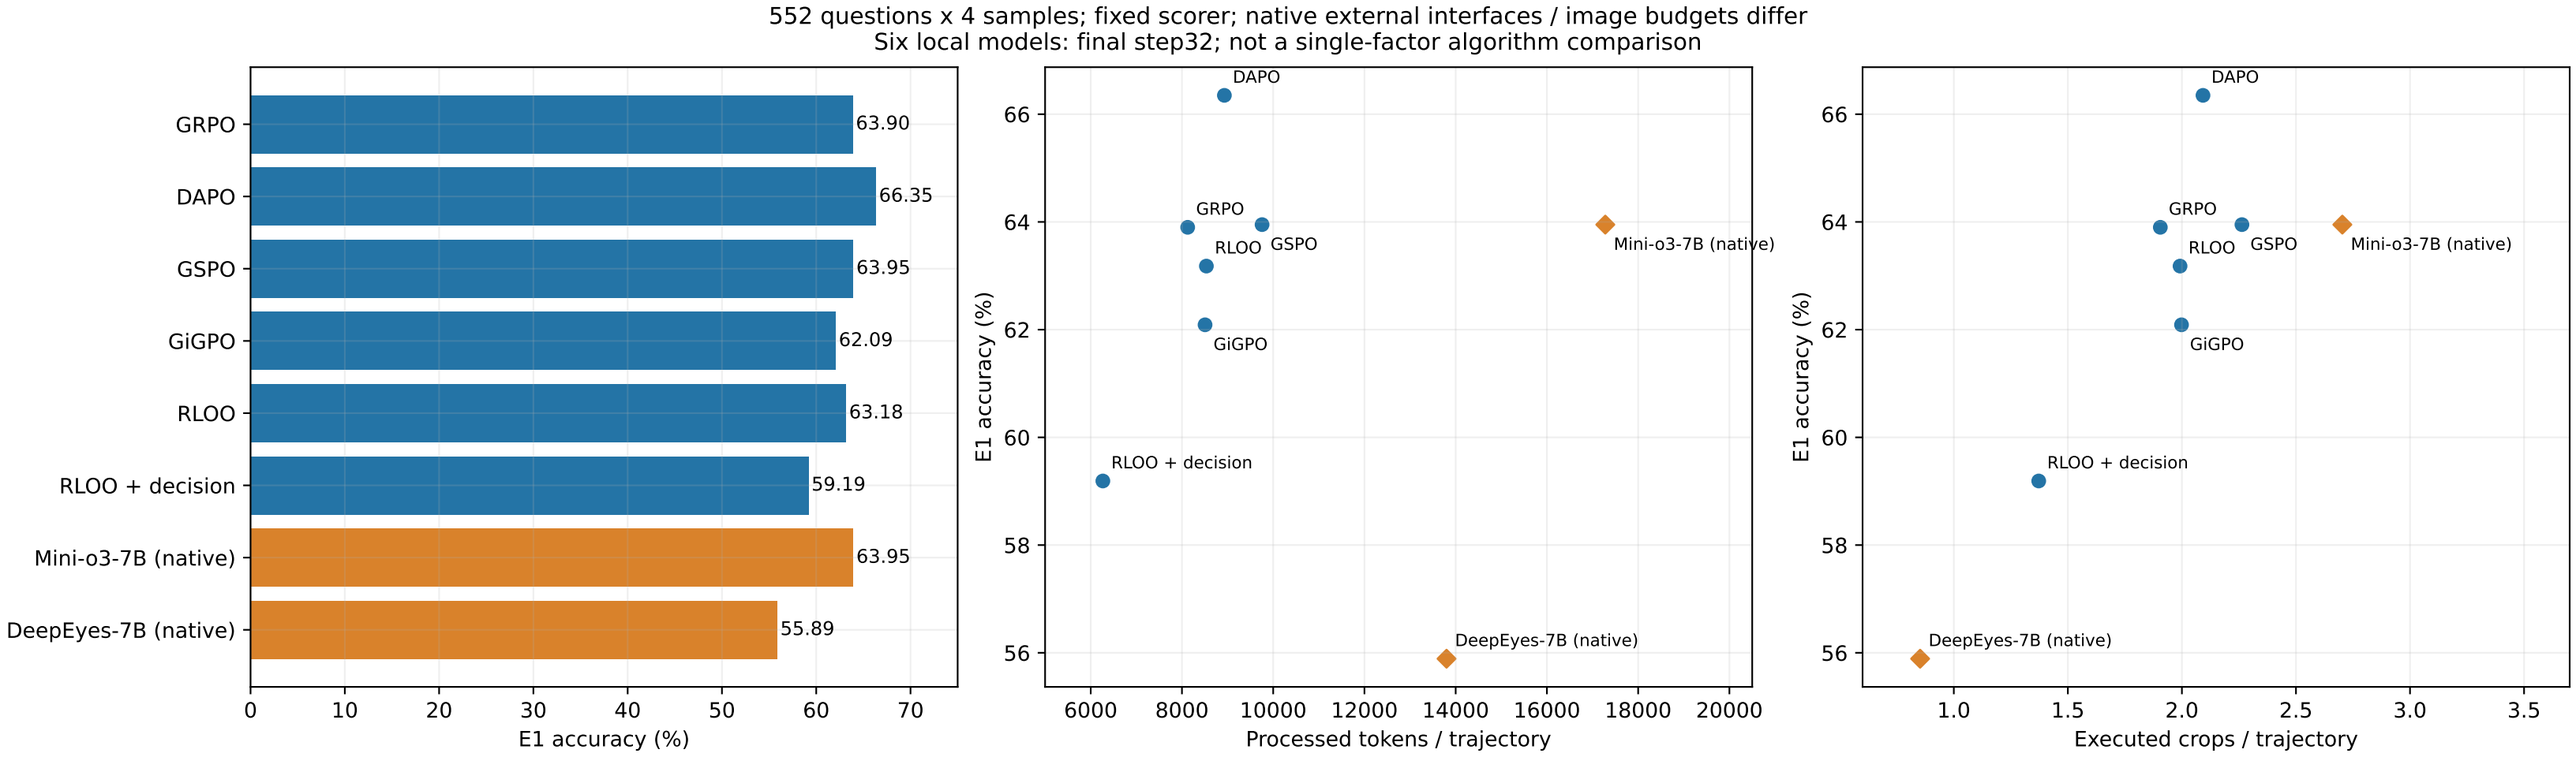
552 questions x 4 samples; fixed scorer; native external interfaces / image budgets differ
Six local models: final step32; not a single-factor algorithm comparison
GRPO
DAPO
GSPO
GiGPO
RLOO
RLOO + decision
Mini-o3-7B (native)
DeepEyes-7B (native)
63.90
66.35
63.95
62.09
63.18
59.19
63.95
55.89
0
10
20
30
40
50
60
70
E1 accuracy (%)
56
58
60
62
64
66
E1 accuracy (%)
6000
8000
10000
12000
14000
16000
18000
20000
Processed tokens / trajectory
1.0
1.5
2.0
2.5
3.0
3.5
Executed crops / trajectory